

# ***Exploratory Data Analysis***

***Data : Red Wine***



**MUJAHID RAZA**



[linkedin.com/in/mujahidraza](https://www.linkedin.com/in/mujahidraza)



[github.com/mujahid777](https://github.com/mujahid777)

[mujahid.raza\\_07](#)



[@mujahidR\\_07](#)



# EDA With Red Wine Data

## Information

### Additional Information

The two datasets are related to red and white variants of the Portuguese "Vinho Verde" wine. Due to privacy and logistic issues, only physicochemical (inputs) and sensory (the output) variables are available (e.g. there is no data about grape types, wine brand, wine selling price, etc.).

These datasets can be viewed as classification or regression tasks. The classes are ordered and not balanced (e.g. there are many more normal wines than excellent or poor ones). Outlier detection algorithms could be used to detect the few excellent or poor wines. Also, we are not sure if all input variables are relevant. So it could be interesting to test feature selection methods.

### Attribute Information

Input variables (based on physicochemical tests):

- 1 - fixed acidity
- 2 - volatile acidity
- 3 - citric acid
- 4 - residual sugar
- 5 - chlorides
- 6 - free sulfur dioxide
- 7 - total sulfur dioxide
- 8 - density
- 9 - pH
- 10 - sulphates
- 11 - alcohol

Output variable (based on sensory data):

- 12 - quality (score between 0 and 10)

In [1]:

```
import pandas as pd
df= pd.read_csv('winequality-red.csv')
```

In [2]:

```
df.head()
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphate
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	(
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	(
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	(
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	(
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	(

In [3]:

```
# summery of the dataset
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

In [4]:

```
# Descriptive summary of the dataset  
df.describe()
```

Out[4]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.87492
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.46015
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000



In [5]:

```
# Shape of the dataset  
df.shape
```

Out[5]:

(1599, 12)

In [6]:

```
# List down the columns  
df.columns
```

Out[6]:

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide',  
      'density',  
      'pH', 'sulphates', 'alcohol', 'quality'],  
      dtype='object')
```

In [7]:

```
df['quality'].unique()
```

Out[7]:

array([5, 6, 7, 4, 8, 3], dtype=int64)

In [8]:

```
# Conclusion-- Imbalanced Dataset
```

```
df['quality'].value_counts()
```

Out[8]:

```
5    681
6    638
7    199
4     53
8     18
3     10
Name: quality, dtype: int64
```

In [9]:

```
# Missing Values
```

```
df.isnull().sum()
```

Out[9]:

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density           0
pH               0
sulphates         0
alcohol           0
quality           0
dtype: int64
```

In [10]:

```
# To check duplicate records
```

```
df.duplicated()
```

Out[10]:

```
0    False
1    False
2    False
3    False
4     True
...
1594  False
1595  False
1596   True
1597  False
1598  False
Length: 1599, dtype: bool
```

In [11]:

```
df[df.duplicated()]
```

Out[11]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	su
4	7.4	0.700	0.00	1.90	0.076	11.0	34.0	0.99780	3.51	
11	7.5	0.500	0.36	6.10	0.071	17.0	102.0	0.99780	3.35	
27	7.9	0.430	0.21	1.60	0.106	10.0	37.0	0.99660	3.17	
40	7.3	0.450	0.36	5.90	0.074	12.0	87.0	0.99780	3.33	
65	7.2	0.725	0.05	4.65	0.086	4.0	11.0	0.99620	3.41	
...	...	...	...	...	...	...	...	...	...	...
1563	7.2	0.695	0.13	2.00	0.076	12.0	20.0	0.99546	3.29	
1564	7.2	0.695	0.13	2.00	0.076	12.0	20.0	0.99546	3.29	
1567	7.2	0.695	0.13	2.00	0.076	12.0	20.0	0.99546	3.29	
1581	6.2	0.560	0.09	1.70	0.053	24.0	32.0	0.99402	3.54	
1596	6.3	0.510	0.13	2.30	0.076	29.0	40.0	0.99574	3.42	

240 rows × 12 columns



In [12]:

```
# Remove the duplicate records  
df.drop_duplicates(inplace=True)
```

In [13]:

```
df.shape
```

Out[13]:

(1359, 12)

In [14]:

```
df.corr()
```

Out[14]:

	<b>fixed acidity</b>	<b>volatile acidity</b>	<b>citric acid</b>	<b>residual sugar</b>	<b>chlorides</b>	<b>free sulfur dioxide</b>	<b>total sulfur dioxide</b>
<b>fixed acidity</b>	1.000000	-0.255124	0.667437	0.111025	0.085886	-0.140580	-0.103777
<b>volatile acidity</b>	-0.255124	1.000000	-0.551248	-0.002449	0.055154	-0.020945	0.071701
<b>citric acid</b>	0.667437	-0.551248	1.000000	0.143892	0.210195	-0.048004	0.047358
<b>residual sugar</b>	0.111025	-0.002449	0.143892	1.000000	0.026656	0.160527	0.201038
<b>chlorides</b>	0.085886	0.055154	0.210195	0.026656	1.000000	0.000749	0.045773
<b>free sulfur dioxide</b>	-0.140580	-0.020945	-0.048004	0.160527	0.000749	1.000000	0.667246
<b>total sulfur dioxide</b>	-0.103777	0.071701	0.047358	0.201038	0.045773	0.667246	1.000000
<b>density</b>	0.670195	0.023943	0.357962	0.324522	0.193592	-0.018071	0.078141
<b>pH</b>	-0.686685	0.247111	-0.550310	-0.083143	-0.270893	0.056631	-0.079257
<b>sulphates</b>	0.190269	-0.256948	0.326062	-0.011837	0.394557	0.054126	0.035291
<b>alcohol</b>	-0.061596	-0.197812	0.105108	0.063281	-0.223824	-0.080125	-0.217829
<b>quality</b>	0.119024	-0.395214	0.228057	0.013640	-0.130988	-0.050463	-0.177855

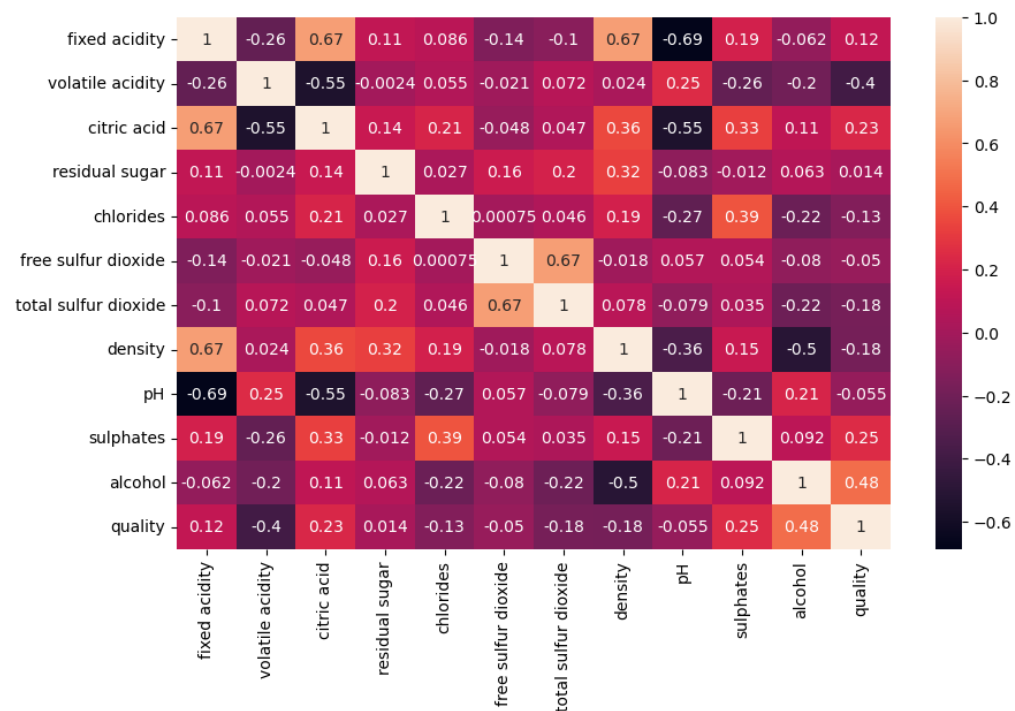


In [15]:

```
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(),annot=True)
```

Out[15]:

<Axes: >



In [16]:

```
df.quality.value_counts()
```

Out[16]:

```
5    577
6    535
7    167
4     53
8     17
3     10
Name: quality, dtype: int64
```




In [17]:

```
df.quality.value_counts().plot(kind='bar')
```

Out[17]:

<Axes: >




In [18]:

```
sns.displot(df['fixed acidity'])
```

Out[18]:

<seaborn.axisgrid.FacetGrid at 0x178a037f970>




In [19]:

```
sns.histplot(df['fixed acidity'],kde=True)
```


Out[19]:

<Axes: xlabel='fixed acidity', ylabel='Count'>



In [20]:

```
for i in df.columns:  
    sns.histplot(df[i],kde=True)
```




In [21]:

```
# Categorical plot
```

```
sns.catplot(x= 'quality',y= 'alcohol',data=df,kind='box')
```

Out[21]:

<seaborn.axisgrid.FacetGrid at 0x178a015ee30>




In [22]:

```
sns.scatterplot(x='alcohol', y='pH', hue='quality', data=df)
```

Out[22]:

<Axes: xlabel='alcohol', ylabel='pH'>




In [23]:

```
sns.catplot(x= 'quality',y= 'citric acid',data=df,kind='box')
```

Out[23]:

<seaborn.axisgrid.FacetGrid at 0x1789cdf3550>




In [24]:

```
sns.catplot(x= 'quality',y= 'free sulfur dioxide',data=df,kind='box')
```

Out[24]:

<seaborn.axisgrid.FacetGrid at 0x178a06b10c0>




In [26]:

```
sns.catplot(x= 'quality',y= 'residual sugar',data=df,kind='box')
```

Out[26]:

<seaborn.axisgrid.FacetGrid at 0x178a26d52d0>



In [ ]: